

DIGITAL SERVICES PROGRAM

AI Agents Bootcamp

Vibe Coding — Websites & AI Agents, No Code Required

COURSE DETAILS & 7-DAY CURRICULUM

Duration	Fee	Certificates	New Batch
7 Days	PKR 10,000	4 (3 Anthropic + 1 DSP)	Every Monday

What You'll Build

Across 7 days you'll go from idea to a live, tested AI agent — and, where applicable, a deployed website front end — using only Claude.ai, the Anthropic Console, and Claude Code/Cowork. No Python, no API keys, no prior coding background needed.

Core Formula

AGENT = CLAUDE + JOB DESCRIPTION + TOOLS + LOOP

Curriculum at a Glance

Day 1	Idea → Agent — mindset, tools stack, first website, agent formula
Day 2	Job Description & Prompting — the 7-part JD
Day 3	Knowledge & Memory — the Memory Ladder
Day 4	Tools & Claude Code — giving agents real actions
Day 5	Test / Secure / Ship — the 10-Question Test Sheet, going live
Sat–Sun	Capstone — build, deploy, and showcase a complete agent

DAY 1

Idea → Agent: Vibe Coding Basics

Mindset Shift

Vibe coding means describing outcomes in plain English, not writing syntax. You direct, Claude builds.

Tools Stack (Zero Setup)

Claude.ai	Chat, planning, prompting, writing Job Descriptions
Claude Artifacts	Instant websites/apps rendered live in-browser
Anthropic Console	Where agents actually run — Workbench for testing
Claude Code / Cowork	File structure, deployment, multi-step builds

Building a Website (Day 1 Version)

1. Prompt Claude with purpose, audience, sections needed, tone, brand colors
2. Ask for a single-file HTML/React artifact first — see it live immediately
3. Iterate by pointing at what's wrong — never touch code yourself
4. Once approved, move to Claude Code/Cowork for real deployment

Core Agent Formula

AGENT = CLAUDE + JOB DESCRIPTION + TOOLS + LOOP

Job Descriptions & Prompting

Why the JD Is Everything

The JD is the contract: it tells Claude who it is, what it's trying to do, and how to behave when things go wrong. Vague JD = inconsistent agent.

The 7-Part Job Description

Role	Who is this agent?
Goal	What outcome it's driving toward, in one sentence
Audience	Who it's talking to
Tone	Formal, friendly, concise, playful — pick one, be explicit
Steps	The exact sequence it should follow, in order
Rules	Hard boundaries — what it must never do or say
Examples	2-3 sample exchanges showing the behavior you want

Core Prompting Principles

- Be specific, not clever
- Show, don't just tell — one good example beats three lines of instruction
- Break big asks into ordered steps
- Set the failure behavior explicitly
- Iterate in small edits, one change at a time

Today's Formula

GOOD JD = CLEAR ROLE + ONE GOAL + NAMED RULES + REAL EXAMPLES

Knowledge & Memory

Two Different Problems

Knowledge	Facts the agent needs to answer correctly — prices, policies, FAQs
Memory	What the agent recalls about this conversation or user, over time

The Memory Ladder

- Rung 1 — None: stateless, single reply, no history
- Rung 2 — Session: remembers within one conversation only
- Rung 3 — User profile: remembers a returning user across sessions
- Rung 4 — Shared knowledge base: pulls from a shared documents/FAQ store

Where Knowledge Comes From

- Paste short, stable facts directly into the JD
- Upload docs/FAQs as reference material for longer content
- Never let it guess — an honest "I don't know" beats a confident wrong answer
- Keep source material updated

Today's Formula

RELIABLE AGENT = RIGHT KNOWLEDGE + RIGHT MEMORY RUNG + HONEST FALLBACK

Tools & Claude Code

What a Tool Actually Is

A defined action the agent may call, with a name, a when-to-use description, and required inputs. Claude decides WHEN to call it; you decide WHICH tools exist.

Common Tool Categories

Search / lookup	Web search, database query, FAQ retrieval — low risk
Scheduling	Calendar check, book a slot, send a reminder
Communication	Send email, post a message, notify a channel
Code execution	Claude Code — run scripts, build files, deploy, fix bugs
Transactions	Payments, bookings — highest risk, needs guardrails

Where Claude Code Fits

- Use once the agent needs real files, a real repo, or deployment
- Still no manual coding — you describe, Claude Code writes and runs it
- Good for turning an Artifact into a deployed site or wiring a tool's backend
- Always review before shipping — you approve, it executes

Today's Formula

USEFUL AGENT = RIGHT TOOL + CLEAR WHEN-TO-USE RULE + SAFE DEFAULTS

Test, Secure, Ship

Test

- Run the DSP 10-Question Test Sheet — happy path, edge cases, hostile inputs
- Try to break the JD with contradictions and pressure to bend a Rule
- Test the tool trigger both ways — should and shouldn't call it
- Have someone else test it who didn't build it

Secure

Prompt injection	Untrusted text is never treated as instructions
Data exposure	Never repeat back private data unasked
Tool guardrails	High-risk actions need explicit confirmation
Rate & scope limits	Cap tool calls and spend per conversation

Ship

- Move from Claude.ai testing to your deployed environment
- Use the One Agent Many Clients model for multi-customer agents
- Deploy the front end via Claude Code/Cowork if website-attached
- Write a one-page handoff document

Today's Formula

SHIP-READY AGENT = PASSED TEST SHEET + CLOSED SECURITY GAPS + LIVE DEPLOYMENT

SAT–SUN

Capstone: Build, Deploy, Showcase

Saturday — Build

- Lock scope first — one agent, one use case
- Finalize the JD end-to-end
- Wire up knowledge and only the tools genuinely needed
- Build the front end if website-attached
- Run the 10-Question Test Sheet before end of day

Sunday — Deploy & Polish

- Fix everything the test sheet caught
- Deploy live
- Write the one-page handoff
- Prepare the showcase: problem → agent → live demo → next steps

Showcase Criteria

Clarity	Can a stranger understand it in 30 seconds?
Reliability	Does it hold up live, not just in a rehearsed path?
Craft	Is the JD tight, tone consistent, front end clean?
Judgment	Was it scoped sensibly, not over-built?

Capstone Formula

SHOWCASE-READY = TIGHT SCOPE + LIVE DEPLOY + CLEAR STORY

After This Weekend

You'll receive your 4 certificates (3 Anthropic + 1 DSP). To go from builder to professional operator — client discovery, requirements gathering, real engagements — the next step is the 1-Month FDE (Forward Deployed Engineer) track.